


# **Advanced Pandas Techniques for Machine Learning**

**With Code Example**

---



# Introduction to Advanced Data Manipulation

Data manipulation is a crucial skill in machine learning and data science. This presentation covers advanced techniques using pandas and Python, focusing on reshaping, transforming, and analyzing complex datasets. We'll explore methods like stacking, unstacking, working with MultiIndex DataFrames, and converting between long and wide data formats.

```
import pandas as pd
import numpy as np

# Create a sample DataFrame
df = pd.DataFrame({
    'Category': ['A', 'A', 'B', 'B'],
    'Year': [2022, 2023, 2022, 2023],
    'Sales': [100, 120, 90, 110],
    'Profit': [20, 25, 18, 22]
})

print(df)
```

Output:

	Category	Year	Sales	Profit
0	A	2022	100	20
1	A	2023	120	25
2	B	2022	90	18
3	B	2023	110	22

# Stacking DataFrames

Stacking is the process of transforming columns into rows, creating a hierarchical index. This technique is useful when you want to reshape your data for analysis or visualization.

```
# Stack the 'Sales' and 'Profit' columns
stacked_df = df.set_index(['Category', 'Year']).stack()

print(stacked_df)
```

Output:

```
Category  Year  Sales  Profit
A         2022   100     20
          2023   120     25
B         2022    90     18
          2023   110     22
dtype: int64
```

# Unstacking DataFrames

Unstacking is the reverse operation of stacking, transforming rows into columns. This can be useful for creating pivot tables or reorganizing data for easier analysis.

```
# Unstack the previously stacked DataFrame  
unstacked_df = stacked_df.unstack(level=-1)  
  
print(unstacked_df)
```

Output:

	Profit	Sales
Category Year		
A 2022	20	100
A 2023	25	120
B 2022	18	90
B 2023	22	110



follow for more

# MultiIndex DataFrames

MultiIndex (or hierarchical index) DataFrames allow you to work with higher-dimensional data in a two-dimensional structure. They are particularly useful for representing complex datasets with multiple levels of categorization.

```
# Create a MultiIndex DataFrame
multi_index_df = pd.DataFrame({
    ('Sales', 'Q1'): [100, 90],
    ('Sales', 'Q2'): [110, 95],
    ('Profit', 'Q1'): [20, 18],
    ('Profit', 'Q2'): [22, 19]
}, index=pd.MultiIndex.from_product([['A', 'B'], [2023]], names=['Category',
'Year']))

print(multi_index_df)
```

Output:

		Sales		Profit	
		Q1	Q2	Q1	Q2
Category	Year				
A	2023	100	110	20	22
B	2023	90	95	18	19

Swipe next →

# Accessing Data in MultiIndex DataFrames

Working with MultiIndex DataFrames requires understanding how to access and manipulate data across different levels of the index hierarchy.

```
# Accessing data using MultiIndex
print(multi_index_df.loc['A', 2023, 'Sales'])
print(multi_index_df.loc[('A', 2023), ('Sales', 'Q1')])

# Selecting specific levels
print(multi_index_df.xs('Q1', axis=1, level=1))
```

Output:

```
Q1    100
Q2    110
Name: (A, 2023, Sales), dtype: int64

100

   Sales  Profit
Category Year
A      2023    100    20
B      2023     90    18
```

# Long vs. Wide Data Formats

follow for more

Data can be represented in long (narrow) or wide formats. Long format is often preferred for analysis and modeling, while wide format can be more readable for humans.

```
# Wide format (original DataFrame)
print("Wide format:")
print(df)

# Convert to long format
long_df = df.melt(id_vars=['Category', 'Year'], var_name='Metric',
                  value_name='Value')
print("\nLong format:")
print(long_df)
```

Output:

```
Wide format:
  Category  Year  Sales  Profit
0        A  2022    100     20
1        A  2023    120     25
2        B  2022     90     18
3        B  2023    110     22

Long format:
  Category  Year  Metric  Value
0        A  2022   Sales    100
1        A  2023   Sales    120
2        B  2022   Sales     90
3        B  2023   Sales    110
4        A  2022  Profit     20
5        A  2023  Profit     25
6        B  2022  Profit     18
7        B  2023  Profit     22
```

Swipe next →

# Melting DataFrames

save for later 

The melt function is used to transform wide-format data into long-format data. This is particularly useful when preparing data for analysis or visualization that requires data in a "tidy" format.

```
# Create a wide-format DataFrame
wide_df = pd.DataFrame({
    'Name': ['Alice', 'Bob'],
    'Math': [90, 85],
    'Science': [88, 92],
    'History': [78, 89]
})

# Melt the DataFrame
melted_df = wide_df.melt(id_vars=['Name'], var_name='Subject', value_name='Score')

print("Original wide format:")
print(wide_df)
print("\nMelted long format:")
print(melted_df)
```

Output:

```
Original wide format:
   Name  Math  Science  History
0  Alice    90      88      78
1   Bob    85      92      89

Melted long format:
   Name  Subject  Score
0  Alice     Math    90
1   Bob     Math    85
2  Alice  Science    88
3   Bob  Science    92
4  Alice  History    78
5   Bob  History    89
```

Swipe next →



# Advanced Grouping and Aggregation

save for later



Pandas provides powerful tools for grouping and aggregating data, allowing you to perform complex analyses on your datasets.

```
# Create a sample DataFrame
df = pd.DataFrame({
    'Category': ['A', 'A', 'B', 'B', 'A', 'B'],
    'Subcategory': ['X', 'Y', 'X', 'Y', 'X', 'Y'],
    'Value': [10, 15, 20, 25, 30, 35]
})

# Perform advanced grouping and aggregation
result = df.groupby(['Category', 'Subcategory']).agg({
    'Value': ['sum', 'mean', 'max']
}).reset_index()

print("Original data:")
print(df)
print("\nGrouped and aggregated data:")
print(result)
```

Output:

```
Original data:
   Category Subcategory  Value
0         A           X     10
1         A           Y     15
2         B           X     20
3         B           Y     25
4         A           X     30
5         B           Y     35

Grouped and aggregated data:
   Category Subcategory  Value
0         A           X     40
1         A           Y     15
2         B           X     20
3         B           Y     60
```

	Category	Subcategory	Value
0	A	X	40
1	A	Y	15
2	B	X	20
3	B	Y	60

Swipe next →

# Handling Missing Data follow for more

Dealing with missing data is a common task in data manipulation. Pandas offers various methods to handle missing values effectively.

```
# Create a DataFrame with missing values
df = pd.DataFrame({
    'A': [1, 2, np.nan, 4],
    'B': [5, np.nan, np.nan, 8],
    'C': [9, 10, 11, 12]
})

print("Original DataFrame:")
print(df)

# Fill missing values
df_filled = df.fillna(method='ffill')
print("\nDataFrame after forward fill:")
print(df_filled)

# Drop rows with any missing values
df_dropped = df.dropna()
print("\nDataFrame after dropping rows with missing values:")
print(df_dropped)
```

Output:

```
Original DataFrame:
   A    B    C
0  1.0  5.0   9
1  2.0  NaN  10
2  NaN  NaN  11
3  4.0  8.0  12

DataFrame after forward fill:
   A    B    C
0  1.0  5.0   9
1  2.0  5.0  10
2  2.0  5.0  11
3  4.0  8.0  12

DataFrame after dropping rows with missing values:
   A    B    C
0  1.0  5.0   9
3  4.0  8.0  12
```

Swipe next →

# Merging and Joining DataFrames

save for later 

Combining data from multiple sources is a common operation in data manipulation. Pandas provides various methods to merge and join DataFrames.

```
# Create two sample DataFrames
df1 = pd.DataFrame({'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']})
df2 = pd.DataFrame({'ID': [2, 3, 4], 'Age': [25, 30, 35]})

# Perform an inner join
inner_join = pd.merge(df1, df2, on='ID', how='inner')

# Perform a left join
left_join = pd.merge(df1, df2, on='ID', how='left')

print("DataFrame 1:")
print(df1)
print("\nDataFrame 2:")
print(df2)
print("\nInner Join:")
print(inner_join)
print("\nLeft Join:")
print(left_join)
```

Output:

```
DataFrame 1:
   ID  Name
0   1  Alice
1   2   Bob
2   3 Charlie

DataFrame 2:
   ID  Age
0   2   25
1   3   30
2   4   35

Inner Join:
   ID  Name  Age
0   2   Bob   25
1   3 Charlie  30

Left Join:
   ID  Name  Age
0   1  Alice  NaN
1   2   Bob  25.0
2   3 Charlie 30.0
```

Swipe next →

# Time Series Data Manipulation

follow for more

Working with time series data is common in many data analysis tasks. Pandas provides powerful tools for manipulating and analyzing time-based data.

```
# Create a time series DataFrame
dates = pd.date_range(start='2023-01-01', end='2023-01-10', freq='D')
ts_df = pd.DataFrame({'Date': dates, 'Value': np.random.randn(len(dates))})

# Set the Date column as the index
ts_df.set_index('Date', inplace=True)

# Resample the data to weekly frequency
weekly_data = ts_df.resample('W').mean()

print("Original daily data:")
print(ts_df)
print("\nResampled weekly data:")
print(weekly_data)
```

Output:

```
Original daily data:
              Value
Date
2023-01-01    0.354314
2023-01-02   -0.499253
2023-01-03    1.209451
2023-01-04    0.443120
2023-01-05   -0.509704
2023-01-06    0.332290
2023-01-07   -1.213095
2023-01-08   -0.278596
2023-01-09    0.332784
2023-01-10    0.580333

Resampled weekly data:
              Value
Date
2023-01-08    0.017018
2023-01-15    0.211507
```

Swipe next →



# Real-Life Example: Weather Data Analysis

save for later 

Let's analyze a weather dataset to demonstrate the application of advanced data manipulation techniques in a real-world scenario.

```
# Create a sample weather dataset
dates = pd.date_range(start='2023-01-01', end='2023-12-31', freq='D')
weather_data = pd.DataFrame({
    'Date': dates,
    'Temperature': np.random.uniform(0, 30, len(dates)),
    'Humidity': np.random.uniform(30, 90, len(dates)),
    'Precipitation': np.random.uniform(0, 50, len(dates))
})

# Set Date as index
weather_data.set_index('Date', inplace=True)

# Calculate monthly averages
monthly_avg = weather_data.resample('M').mean()

# Find the hottest day of each month
hottest_days = weather_data.resample('M')['Temperature'].idxmax()

print("Monthly averages:")
print(monthly_avg.head())
print("\nHottest days of each month:")
print(hottest_days.head())
```

Output:

```
Monthly averages:
              Temperature  Humidity  Precipitation
Date
2023-01-31      14.819144    58.741148      24.776924
2023-02-28      15.326223    60.853798      23.865147
2023-03-31      14.606469    59.707280      24.705367
2023-04-30      15.402669    60.167848      25.380250
2023-05-31      15.066451    59.729830      24.911770

Hottest days of each month:
Date
2023-01-31    2023-01-24
2023-02-28    2023-02-25
2023-03-31    2023-03-29
2023-04-30    2023-04-29
2023-05-31    2023-05-27
Freq: M, Name: Temperature, dtype: datetime64[ns]
```

Swipe next →

# Real-Life Example: Customer Segmentation

In this example, we'll demonstrate how to use advanced data manipulation techniques for customer segmentation based on purchasing behavior.

```
import pandas as pd
import numpy as np

# Create a sample customer purchase dataset
np.random.seed(42)
n_customers = 1000
customer_data = pd.DataFrame({
    'CustomerID': range(1, n_customers + 1),
    'TotalPurchases': np.random.randint(1, 100, n_customers),
    'AvgOrderValue': np.random.uniform(10, 200, n_customers),
    'DaysSinceLastPurchase': np.random.randint(1, 365, n_customers)
})

# Perform customer segmentation
customer_data['ValueSegment'] = pd.qcut(customer_data['AvgOrderValue'], q=3,
labels=['Low', 'Medium', 'High'])
customer_data['FrequencySegment'] = pd.qcut(customer_data['TotalPurchases'], q=3,
labels=['Low', 'Medium', 'High'])
customer_data['RecencySegment'] = pd.qcut(customer_data['DaysSinceLastPurchase'],
q=3, labels=['High', 'Medium', 'Low'])

# Analyze segments
segment_analysis = customer_data.groupby(['ValueSegment', 'FrequencySegment',
'RecencySegment']).size().unstack(fill_value=0)

print("Sample of customer data:")
print(customer_data.head())
print("\nSegment analysis:")
print(segment_analysis)
```

# Output

save for later 

Sample of customer data:

CustomerID	TotalPurchases	AvgOrderValue	DaysSinceLastPurchase	ValueSegment	
FrequencySegment	RecencySegment				
0	1	37	124.472172	143	High
Medium	Medium				
1	2	87	118.256325	112	High
High	Medium				
2	3	52	192.379052	343	High
High	Low				
3	4	69	82.804886	134	Medium
High	Medium				
4	5	66	155.693686	335	High
High	Low				

Segment analysis:

ValueSegment	FrequencySegment	RecencySegment	High	Low	Medium
High	High	High	29	23	32
		Low	39	25	35
		Medium	38	30	32
	Low	High	5	25	18
		Low	9	24	25
		Medium	8	23	20
	Medium	High	16	21	30
		Low	21	24	26
		Medium	19	24	24
Medium	High	High	31	19	33
		Low	34	23	32
		Medium	29	28	26
	Low	High	8	24	22
		Low	7	25	22
		Medium	11	18	24
	Medium	High	14	27	25
		Low	21	20	24
		Medium	22	19	26
Low	High	High	24	27	26
		Low	31	23	30
		Medium	30	24	29
	Low	High	9	17	24
		Low	9	25	23
		Medium	8	26	21
	Medium	High	16	25	26
		Low	16	22	26
		Medium	17	24	26

Swipe next →

# Additional Resources

For further exploration of advanced data manipulation techniques using pandas and Python, consider the following resources:

1. pandas documentation:  
<https://pandas.pydata.org/docs/>
2. "Python for Data Analysis" by Wes McKinney (creator of pandas)
3. DataCamp's "Data Manipulation with pandas" course
4. Real Python's pandas tutorials:  
<https://realpython.com/learning-paths/pandas-data-science/>
5. ArXiv paper on data manipulation techniques: "A Survey on Data Preprocessing for Data Mining" (<https://arxiv.org/abs/1811.05242>)



# **Data scientist & ML Engineer**



**Follow For More Data  
Science Content**