

Simple Roadmap

to start your career in
Data Science



Chapter 1: Foundation:

The first and most important thing to learn for data science is mathematics. It forms the basis of many data science concepts and techniques. Mathematics helps you understand how data is represented and transformed, how models learn and make predictions, and how to analyze and interpret data to make informed decisions.

Mathematics for Data Science constitutes a few key skills you need to develop:

Linear Algebra:

- **Key Concepts:** Vectors, matrices, linear transformations, eigenvalues, and eigenvectors.
- **Why it's important:** Linear algebra forms the foundation of how data is represented and transformed in machine learning models. It's crucial for tasks like dimensionality reduction, image processing, and natural language processing.
- **Resources and Courses:**
 - **3Blue1Brown's Essence of Linear Algebra series on YouTube:** This is a fantastic visual introduction to linear algebra concepts.
 - **Khan Academy's Linear Algebra course:** Offers a comprehensive and beginner-friendly approach to learning linear algebra.
 - **"Mathematics for Machine Learning: Linear Algebra"** on Coursera: This course combines theory with hands-on Python exercises, allowing you to apply linear algebra concepts in a practical setting.

Calculus:

- **Key Concepts:** Derivatives, integrals, gradients, optimization algorithms.
- **Why it's important:** Calculus helps you understand how models learn from data, fine-tune their parameters, and make accurate predictions.
- **Resources and Courses:**
 - Khan Academy's Calculus courses: These provide a solid foundation in calculus concepts, starting from the basics.
 - "Calculus 1A: Differentiation" on edX: This course focuses on the applications of calculus, particularly differentiation, in various fields.
 - MIT OpenCourseWare's Single Variable Calculus: Offers a more rigorous and in-depth exploration of calculus.

Statistics:

- **Key Concepts:** Descriptive statistics, probability distributions, hypothesis testing, regression analysis.
- **Why it's important:** Statistics is essential for understanding data, evaluating model performance, and making informed decisions based on evidence.
- **Resources and Courses:**
 - "Statistics with Python Specialization" on Coursera: This specialization covers a wide range of statistical concepts and their implementation in Python.
 - "Data Analysis and Statistical Inference" on edX: This course focuses on applying statistics to real-world datasets and drawing meaningful conclusions.
 - Khan Academy's Statistics and Probability course: Provides a beginner-friendly introduction to statistics and probability.

Chapter 2: Programming

when it comes to data science, Python is your trusty Swiss Army knife. It's versatile, easy to learn, and has a massive community and ecosystem of libraries built specifically for data science tasks.

Why Python?

Python's simplicity and readability make it a great choice for beginners. You can focus on learning data science concepts without getting bogged down by complex syntax. Plus, its vast collection of libraries provides pre-built tools for everything from data cleaning and manipulation to machine learning and visualization.

Key Concepts to Master

- **Basics:** Variables, data types, control flow (if-else statements, loops), functions, and object-oriented programming. These fundamentals form the building blocks of any Python program.
- **Data Structures:** Lists, tuples, dictionaries, and sets. These are essential for organizing and storing data in a way that's easy to work with.
- **File Handling:** Reading and writing data from various file formats (CSV, Excel, JSON, etc.). This is crucial for importing and exporting data for analysis.
- **Exception Handling:** Dealing with errors gracefully to prevent your programs from crashing unexpectedly.

Resources and Courses

Interactive Learning:

- Codecademy's "Learn Python 3": This interactive course is a fantastic way to get hands-on with Python basics. It guides you through exercises and projects, making learning fun and engaging.
- DataCamp's "Introduction to Python": Another excellent interactive platform with courses specifically tailored for data science.

Comprehensive Courses:

- "Python for Everybody Specialization" on Coursera: This specialization provides a solid foundation in Python programming and introduces you to data analysis using Python.
- "Complete Python Bootcamp: Go from zero to hero in Python 3" on Udemy: This popular course covers a wide range of Python topics, from the basics to more advanced concepts.

Actionable Steps:

1. **Install Python:** If you haven't already, download and install the latest version of Python from the official website (<https://www.python.org/>).
2. **Choose a learning resource:** Select a course or platform that suits your learning style and preferences.
3. **Set a schedule:** Dedicate regular time each week to learning and practicing Python. Consistency is key!
4. **Practice, practice, practice:** The more you code, the better you'll become. Start with simple exercises and gradually work your way up to more complex projects.
5. **Explore libraries:** Once you have a grasp of the basics, start exploring essential data science libraries like NumPy, Pandas, Matplotlib, and Seaborn. We'll cover these in more detail in the next chapter.

Chapter 3: Libraries:

Now let's explore the powerful libraries that will supercharge your data science capabilities. Think of these libraries as your trusty sidekicks, ready to help you tackle any data challenge that comes your way.

NumPy: The Numerical Ninja

- **What it does:** NumPy is the foundation for numerical computing in Python. It provides powerful tools for working with arrays and matrices, performing mathematical operations, and handling large datasets efficiently.
- **Why it's important:** NumPy's arrays are much faster and more memory-efficient than Python's built-in lists, making it essential for handling the large datasets commonly encountered in data science.
 - Key features: N-dimensional arrays (ndarrays)
 - Mathematical operations on arrays (element-wise, matrix multiplication, etc.)
 - Random number generation
 - Linear algebra routines
 - Fourier transforms

Pandas: The Data Wrangler

- **What it does:** Pandas is your go-to library for data manipulation and analysis. It introduces the DataFrame, a powerful data structure similar to a spreadsheet but with much more functionality.
- **Why it's important:** Pandas makes it easy to clean, transform, and explore your data, preparing it for further analysis and machine learning.
 - Key features: DataFrames and Series (1-dimensional arrays)
 - Reading and writing data from various file formats (CSV, Excel, SQL databases, etc.)
 - Handling missing data
 - Filtering, sorting, and grouping data
 - Merging and joining datasets
 - Time series analysis

Matplotlib & Seaborn: The Visualization

- **What they do:** Matplotlib is the core plotting library in Python, providing a wide range of customizable plotting options. Seaborn builds on top of Matplotlib, offering a higher-level interface for creating statistically meaningful and aesthetically pleasing visualizations.
- **Why they're important:** Visualizing data is crucial for understanding patterns, trends, and outliers. It helps you communicate your findings effectively and make data-driven decisions.
 - Key features: Line plots, scatter plots, bar charts, histograms, heatmaps, and more
 - Customization options for colors, labels, legends, and annotations
 - Seaborn's built-in statistical plotting functions for visualizing distributions, relationships, and comparisons

Scikit-learn: The Machine Learning

- **What it does:** Scikit-learn is your one-stop shop for machine learning in Python. It provides a comprehensive collection of algorithms and tools for tasks like classification, regression, clustering, and dimensionality reduction.
- **Why it's important:** Scikit-learn makes it easy to build, train, and evaluate machine learning models, enabling you to extract insights and make predictions from your data.
 - Key features: Supervised learning algorithms (linear regression, logistic regression, decision trees, random forests, support vector machines, etc.)
 - Unsupervised learning algorithms (k-means clustering, hierarchical clustering, principal component analysis, etc.)
 - Model selection and evaluation tools
 - Preprocessing and feature engineering techniques

Actionable Steps:

1. **Install and Explore:** If you haven't already, install these libraries using **`pip install numpy pandas matplotlib seaborn scikit-learn`**. Then, start exploring their basic functionalities by following tutorials and examples.
2. **Learn by Doing:** The best way to master these libraries is to use them in real projects. Start with simple datasets and gradually work your way up to more complex challenges.
3. **Course Recommendation:** "Python for Data Science and AI" on Coursera is an excellent resource for learning how to use these libraries effectively for data analysis and machine learning.

Chapter 4: Data Wrangling and Visualization

Now that you've got your Python and library skills sharpened, it's time to tackle the often messy and unruly world of data. In this chapter, we'll explore the crucial skills of data wrangling and visualization, which will empower you to transform raw data into valuable insights.

Data Wrangling: Cleaning and Preparing Your Data

Think of data wrangling as the essential housekeeping before the main event. It involves cleaning, transforming, and organizing your data to make it suitable for analysis and machine learning. This step is often the most time-consuming but also the most critical in any data science project.

Key tasks in data wrangling:

- **Handling missing values:** Identifying and dealing with missing data points.
- **Dealing with outliers:** Detecting and addressing extreme or unusual values that can skew your analysis.
- **Data cleaning:** Fixing errors, inconsistencies, and formatting issues in your data.
- **Data transformation:** Converting data into a suitable format for analysis (e.g., converting categorical variables to numerical representations).
- **Feature engineering:** Creating new features or transforming existing ones to improve the performance of your machine learning models.

Visualization: Painting a Picture with Data

Data visualization is the art of communicating insights through charts, graphs, and other visual representations. It helps you understand patterns, trends, and outliers in your data, and it's a powerful way to share your findings with others.

- **Key types of visualizations:**
 - Bar charts and histograms: Visualize the distribution of categorical or numerical data.
 - Scatter plots: Explore relationships between two numerical variables.
 - Line charts: Track changes over time.
 - Heatmaps: Visualize the intensity of values across two dimensions.
 - And many more!

Tools and Techniques

- **Pandas:** Your primary tool for data wrangling and basic visualization. Use it to load, clean, transform, and explore your data.
- **Matplotlib & Seaborn:** Your go-to libraries for creating a wide range of visualizations. Matplotlib offers flexibility and customization, while Seaborn provides a higher-level interface for statistically meaningful plots.
- **Other libraries:** Consider exploring additional libraries like Plotly and Bokeh for interactive visualizations, and Folium for creating maps.

Actionable Steps:

1. **Get your hands dirty:** Find a dataset that interests you and start exploring it using Pandas. Practice cleaning, transforming, and visualizing the data to uncover insights.
2. **Learn by example:** Follow tutorials and online courses that demonstrate data wrangling and visualization techniques.
3. **Experiment with different visualizations:** Try out various chart types and see how they can reveal different aspects of your data.
4. **Tell a story with your visualizations:** Use your visualizations to communicate a clear and compelling narrative about your findings.

Course Recommendations:

- **Data Visualization with Python on Coursera:** A comprehensive course covering data visualization principles and techniques using Matplotlib and Seaborn.
- **Data Cleaning in Python on DataCamp:** Learn how to handle missing data, outliers, and other data quality issues using Pandas.
- **Kaggle's "Titanic: Machine Learning from Disaster"** competition: A great hands-on project to practice data wrangling and visualization skills.

Chapter 5: Machine Learning:

It's time to dive into the heart of data science: Machine Learning. This is where the magic happens! Machine learning empowers you to build models that can learn from data and make predictions or decisions without being explicitly programmed. It's like teaching a computer to think and learn from experience.

Types of Machine Learning

There are three main types of machine learning:

1. **Supervised Learning:** You provide the model with labeled data, where the input features and the desired output (target) are known. The model learns to map inputs to outputs, allowing it to make predictions on new, unseen data.
 - Common tasks: Classification (predicting categories, e.g., spam vs. not spam) and regression (predicting continuous values, e.g., house prices).
 - Algorithms: Linear regression, logistic regression, decision trees, random forests, support vector machines, neural networks.

Unsupervised learning:

Unsupervised Learning: You provide the model with unlabeled data, and it tries to discover patterns or structures in the data on its own.

- **Common tasks:** Clustering (grouping similar data points together), dimensionality reduction (reducing the number of features while preserving important information), anomaly detection (identifying unusual data points).
- **Algorithms:** K-means clustering, hierarchical clustering, principal component analysis (PCA), t-SNE.
- **Reinforcement Learning:** The model learns to make decisions by interacting with an environment. It receives rewards or penalties based on its actions and learns to maximize its rewards over time.
- **Common tasks:** Game playing, robotics, self-driving cars.
- **Algorithms:** Q-learning, deep Q-networks (DQN), policy gradients.

Key Concepts and Steps:

1. **Problem Formulation:** Clearly define the problem you want to solve and identify the type of machine learning that's best suited for it.
2. **Data Collection and Preparation:** Gather and clean your data, ensuring it's in a suitable format for your chosen algorithms.
3. **Feature Engineering:** Select, transform, or create relevant features that will help your model learn effectively.
4. **Model Selection and Training:** Choose appropriate algorithms, train your models on the data, and tune their parameters to optimize performance.
5. **Model Evaluation:** Assess how well your model generalizes to new, unseen data using appropriate metrics.
6. **Deployment:** Integrate your model into a real-world application or system.

Tools and Libraries:

- **Scikit-learn**: Your primary tool for implementing and evaluating a wide range of machine learning algorithms.
- **TensorFlow & PyTorch**: Powerful deep learning frameworks for building and training complex neural networks.
- **Other libraries**: Explore libraries like XGBoost, LightGBM, and CatBoost for gradient boosting algorithms, and Keras for a user-friendly interface to deep learning.

Actionable Steps:

1. **Start with the basics:** Focus on understanding the core concepts of supervised and unsupervised learning.
2. **Practice with Scikit-learn:** Work through tutorials and examples using Scikit-learn to build and evaluate different models.
3. **Explore Kaggle competitions:** Participate in Kaggle competitions to gain practical experience and learn from others.
4. **Dive into deep learning:** If you're interested in advanced topics, consider exploring deep learning frameworks like TensorFlow or PyTorch.

Course Recommendations:

- **Machine Learning on Coursera by Andrew Ng:** A classic and highly recommended course covering the fundamentals of machine learning.
- **Supervised Machine Learning: Regression and Classification on Coursera:** Focuses on the two most common supervised learning tasks.
- **Deep Learning Specialization on Coursera:** An in-depth exploration of deep learning concepts and techniques.

Chapter 6: Model Evaluation and Selection:

Great job making it this far! Now that you've built and trained several machine learning models, it's time to figure out which one is the champion. In this chapter, we'll explore the crucial steps of model evaluation and selection, helping you choose the best model for your specific task.

Why Model Evaluation Matters

Imagine you're training for a marathon. You wouldn't just run aimlessly without tracking your progress, right? Similarly, in machine learning, you need to evaluate your models to understand how well they're performing and identify areas for improvement.

Model evaluation helps you:

- **Assess performance:** Measure how well your model is doing on various metrics like accuracy, precision, recall, F1-score, and more.
- **Compare models:** Compare the performance of different models to choose the best one for your task.
- **Avoid overfitting:** Ensure your model generalizes well to new, unseen data and doesn't just memorize the training data.
- **Identify areas for improvement:** Pinpoint weaknesses in your model and explore ways to enhance its performance.

Key Evaluation Metrics

The choice of evaluation metrics depends on the type of problem you're solving.

Classification:

- Accuracy: The overall proportion of correct predictions.
- Precision: The proportion of positive predictions that were actually correct.
- Recall: The proportion of actual positives that were correctly identified.
- F1-score: A harmonic mean of precision and recall, providing a balanced measure of performance.
- Confusion matrix: A table that shows the number of true positives, true negatives, false positives, and false negatives.
- ROC curve and AUC: Visualize the trade-off between true positive rate and false positive rate.

Regression:

- **Mean squared error (MSE):** The average squared difference between the predicted values and the actual values.
- **Root mean squared error (RMSE):** The square root of the MSE, providing an error in the same units as the target variable.
- **Mean absolute error (MAE):** The average absolute difference between the predicted values and the actual values.
- **R-squared:** The proportion of variance in the target variable explained by the model.

Model Selection Strategies

- **Train-test split:** Split your data into a training set (used to train the model) and a test set (used to evaluate the model's performance on unseen data).
- **Cross-validation:** Divide your data into multiple folds, train and evaluate your model on different combinations of folds to get a more robust estimate of performance.
- **Hyperparameter tuning:** Experiment with different hyperparameter values (parameters that are not learned from the data) to find the optimal configuration for your model.
- **Grid search and random search:** Systematic approaches to explore different hyperparameter combinations.

Tools & Libraries

- **Scikit-learn:** Provides a wide range of evaluation metrics and tools for model selection and hyperparameter tuning.
- **Other libraries:** Consider exploring libraries like Optuna or Hyperopt for more advanced hyperparameter optimization techniques.

Actionable Steps:

1. **Understand your metrics:** Choose appropriate evaluation metrics based on your problem type and business goals.
2. **Implement train-test split or cross-validation:** Split your data or use cross-validation to get a reliable estimate of your model's performance on unseen data.
3. **Experiment with hyperparameter tuning:** Use grid search or random search to find the best hyperparameters for your model.
4. **Compare and select the best model:** Choose the model that performs best on your chosen evaluation metrics while considering factors like complexity and interpretability.

Tools & Libraries

- **Scikit-learn:** Provides a wide range of evaluation metrics and tools for model selection and hyperparameter tuning.
- **Other libraries:** Consider exploring libraries like Optuna or Hyperopt for more advanced hyperparameter optimization techniques.

Actionable Steps:

1. **Understand your metrics:** Choose appropriate evaluation metrics based on your problem type and business goals.
2. **Implement train-test split or cross-validation:** Split your data or use cross-validation to get a reliable estimate of your model's performance on unseen data.
3. **Experiment with hyperparameter tuning:** Use grid search or random search to find the best hyperparameters for your model.
4. **Compare and select the best model:** Choose the model that performs best on your chosen evaluation metrics while considering factors like complexity and interpretability.

Chapter 7: Deep Learning:

Alright, it's time to take your machine learning skills to the next level with deep learning. Deep learning is a subfield of machine learning that focuses on artificial neural networks, which are inspired by the structure and function of the human brain. These networks can learn complex patterns and representations from data, making them incredibly powerful for tasks like image recognition, natural language processing, and speech recognition.

What are Neural Networks?

Neural networks are composed of interconnected layers of artificial neurons. Each neuron receives input, performs a computation, and passes the output to the next layer. The connections between neurons have weights that are adjusted during training, allowing the network to learn from the data.

Key components of a neural network:

- Input layer: Receives the raw data.
- Hidden layers: Perform computations and learn representations of the data.
- Output layer: Produces the final prediction or decision.
- Weights and biases: Parameters that are adjusted during training to optimize the network's performance.
- Activation functions: Introduce non-linearity into the network, allowing it to learn complex patterns.

Types of Neural Networks

There are many different types of neural networks, each designed for specific tasks. Some common ones include:

- **Feedforward Neural Networks (FNNs)**: The simplest type of neural network, where information flows in one direction from the input layer to the output layer.
- **Convolutional Neural Networks (CNNs)**: Excellent for image recognition and computer vision tasks. They use convolutional layers to extract features from images.
- **Recurrent Neural Networks (RNNs)**: Ideal for sequential data like time series or natural language. They have connections that loop back on themselves, allowing them to maintain memory of past inputs.
- **Long Short-Term Memory (LSTM) networks and Gated Recurrent Units (GRUs)**: Specialized types of RNNs that are better at handling long-term dependencies in sequential data.

Key Concepts and Steps in Deep Learning

1. **Data Preparation:** Gather and preprocess your data, ensuring it's in a suitable format for your chosen neural network architecture.
2. **Network Architecture:** Design the structure of your neural network, including the number of layers, neurons per layer, and activation functions.
3. **Loss Function and Optimizer:** Choose a loss function to measure the error between your model's predictions and the true values, and an optimizer to adjust the weights and biases during training to minimize the loss.
4. **Training:** Feed your data into the network and let it learn by adjusting its weights and biases through backpropagation.
5. **Evaluation and Tuning:** Assess your model's performance on a validation set and fine-tune its hyperparameters to improve its generalization ability.
6. **Deployment:** Integrate your trained model into a real-world application or system.

Tools and Libraries

- **TensorFlow**: A popular and powerful open-source deep learning framework developed by Google.
- **PyTorch**: Another widely used deep learning framework known for its flexibility and dynamic computation graphs.
- **Keras**: A high-level API that simplifies the process of building and training neural networks. It can run on top of TensorFlow, Theano, or CNTK.

Actionable Steps:

1. **Choose a framework:** Start with either TensorFlow or PyTorch, depending on your preference and the specific requirements of your project.
2. **Follow tutorials and examples:** There are numerous online tutorials and examples available for both TensorFlow and PyTorch. Start with simple examples and gradually work your way up to more complex projects.
3. **Practice on real datasets:** Apply your deep learning skills to real-world datasets to gain practical experience and build your portfolio.
4. **Join online communities:** Connect with other deep learning enthusiasts and experts through forums, social media groups, and online communities.

Course Recommendations:

- **Deep Learning Specialization on Coursera:** A comprehensive and highly recommended specialization covering the fundamentals of deep learning and its applications.
- **Practical Deep Learning for Coders by fast.ai:** A hands-on course that teaches you how to build and train deep learning models using PyTorch.
- **Introduction to Deep Learning on MIT OpenCourseWare:** A more theoretical course that delves into the mathematical foundations of deep learning.

Chapter 8: Advanced Topics:

Congratulations on reaching this point in your data science journey! You've built a solid foundation in mathematics, programming, data wrangling, visualization, and machine learning. Now, it's time to expand your horizons and explore some advanced topics that will take your skills to the next level.

Natural Language Processing (NLP)

NLP is a fascinating field that focuses on enabling computers to understand, interpret, and generate human language. It has numerous applications, from chatbots and machine translation to sentiment analysis and text summarization.

- **Key concepts:**

- Text preprocessing: Cleaning and preparing text data for analysis (tokenization, stemming, lemmatization, stop word removal).
- Word embeddings: Representing words as vectors in a high-dimensional space, capturing their semantic meaning and relationships.
- Language models: Powerful models that can generate text, translate languages, and answer questions.
- Sentiment analysis: Determining the emotional tone of a piece of text (positive, negative, or neutral).
- Text summarization: Generating concise summaries of longer documents.

Tools and Libraries:

- **NLTK (Natural Language Toolkit):** A comprehensive library for NLP tasks, providing tools for tokenization, stemming, tagging, parsing, and more.
- **spaCy:** Another popular NLP library known for its speed and efficiency.
- **Gensim:** A library for topic modeling and word embeddings.
- **Transformers:** A state-of-the-art library for building and using transformer-based language models like BERT and GPT.

Actionable Steps:

1. **Start with the basics:** Learn about text preprocessing techniques and how to use NLTK or spaCy for basic NLP tasks.
2. **Explore word embeddings:** Understand how word embeddings work and experiment with pre-trained models like Word2Vec or GloVe.
3. **Dive into language models:** Learn about transformer-based models and how to use them for tasks like text generation and question answering.
4. **Build NLP projects:** Apply your skills to real-world NLP challenges, such as building a sentiment analysis classifier or a text summarization tool.

Course Recommendations:

- **Natural Language Processing Specialization on Coursera:** A comprehensive specialization covering various NLP concepts and techniques.
- **Practical Deep Learning for Coders (NLP) by fast.ai:** A hands-on course that teaches you how to build and train NLP models using PyTorch.

Other Advanced Topics

Besides NLP, there are many other exciting areas you can explore to further enhance your data science expertise:

- **Computer Vision:** Focuses on enabling computers to understand and interpret visual information from the world, such as images and videos.
- **Time Series Analysis:** Analyzes data collected over time to identify trends, seasonality, and other patterns.
- **Recommender Systems:** Build systems that suggest products, movies, or other items to users based on their preferences and behavior.
- **Big Data Technologies:** Learn to work with massive datasets using tools like Apache Spark and Hadoop.
- **Cloud Computing:** Leverage cloud platforms like AWS, Azure, or GCP for scalable data storage and processing.

Other Advanced Topics

Besides NLP, there are many other exciting areas you can explore to further enhance your data science expertise:

- **Computer Vision:** Focuses on enabling computers to understand and interpret visual information from the world, such as images and videos.
- **Time Series Analysis:** Analyzes data collected over time to identify trends, seasonality, and other patterns.
- **Recommender Systems:** Build systems that suggest products, movies, or other items to users based on their preferences and behavior.
- **Big Data Technologies:** Learn to work with massive datasets using tools like Apache Spark and Hadoop.
- **Cloud Computing:** Leverage cloud platforms like AWS, Azure, or GCP for scalable data storage and processing.

Project Ideas to master skills

Exploratory Data Analysis (EDA) on a Dataset of Your Choice:

- Why it's great: EDA is the foundation of any data science project. It helps you understand the data, uncover patterns, and identify potential insights.
- Datasets: Choose a dataset that interests you! Kaggle, UCI Machine Learning Repository, and government websites offer a wealth of options.
- Skills you'll develop: Data cleaning, visualization, statistical analysis, storytelling.

Build a Movie Recommender System:

- Why it's great: Recommender systems are widely used in e-commerce, streaming services, and more. This project is both fun and practical.
- Skills you'll develop: Data preprocessing, collaborative filtering, content-based filtering, model evaluation.

Project Ideas to master skills

Sentiment Analysis on Social Media Data:

- Why it's great: Analyze tweets, reviews, or comments to understand public opinion on a topic or product.
- Skills you'll develop: Text preprocessing, natural language processing (NLP), machine learning classification.

Predict Customer Churn:

- Why it's great: Help businesses identify customers at risk of leaving and take proactive measures to retain them.
- Skills you'll develop: Data cleaning, feature engineering, classification algorithms, model evaluation.

Project Ideas to master skills

Stock Price Prediction:

- Why it's great: A classic project with real-world applications in finance.
- Skills you'll develop: Time series analysis, feature engineering, regression algorithms, model evaluation.

Image Classification with Convolutional Neural Networks (CNNs):

- Why it's great: Build a model that can recognize objects or classify images.
- Skills you'll develop: Deep learning, CNNs, image preprocessing, model training and evaluation.

Project Ideas to master skills

Build a Chatbot:

- Why it's great: Create a conversational AI that can interact with users and provide information or assistance.
- Skills you'll develop: Natural language processing (NLP), dialogue management, machine learning.