

HOW I STARTED LEARNING ML

**From Developer
to AI Research
Scientist**



Bridging Theory and Practice

I consider myself an AI Research Scientist who loves exploring new AI trends and understanding their background. But what truly excites me is turning scientific research into real products. I believe every breakthrough in AI science eventually becomes a solution to real-world problems.

Why this matters:

- Research needs practical application
- Products solve user problems
- Innovation drives development
- Theory enables practice



The Beginning (2016)

Back in 2016, machine learning wasn't as popular as it is today. Everyone talked about it, but few were actually implementing it. As a developer with years of experience, I faced the classic beginner's dilemma - where to start on this complex journey?

Initial challenges:

- Choosing between mathematics and statistics
- Finding the right learning path
- Selecting reliable resources
- Building proper foundations



The Foundation Course

Everything changed when I discovered a comprehensive Stanford machine learning course. It wasn't just another online course - it provided all the necessary calculus and ML basics needed to start. The best part? You don't need to know everything in depth to begin building.

Key learning approach:

- Focus on understanding concepts
- Learn necessary mathematics
- Master core algorithms
- Write everything by hand



Resources



Stanford CS229: Machine Learning Course, Lecture 1 - Andrew Ng (Autumn 2018)



Stanford Online 660K subscribers

Subscribe

40K



1



Share



Download



Save



 DIVE INTO
DEEP LEARNING

D2L.ai: Interactive Deep Learning Book with Multi-Framework Code, Math, and Discussions

 Continuous Integration failing

[Book website](#) | [STAT 157 Course at UC Berkeley](#)

The best way to understand deep learning is learning by doing.

dden units and, for a given time step t , the mini-
mum of inputs: d) and the hidden state of the last
 $\mathbb{R}^{n \times h}$, forget gate $F_t \in \mathbb{R}^{n \times h}$, and output gate

$$I_t = \sigma(X_t W_{xi} + H_{t-1} W_{hi} + b_i),$$
$$F_t = \sigma(X_t W_{xf} + H_{t-1} W_{hf} + b_f),$$
$$O_t = \sigma(X_t W_{xo} + H_{t-1} W_{ho} + b_o),$$

and $W_{hi}, W_{hf}, W_{ho} \in \mathbb{R}^{n \times h}$ are weight param

```
def lstm(inputs, state, params):
    W_xi, W_hi, b_i, W_xf, b_f, W_xo, b_o, W_ho, b_h = params
    (H, C) = state
    outputs = []
    for X in inputs:
        I = nd.sigmoid(nd.dot(X, W_xi) + nd.dot(H, W_hi) + b_i)
        F = nd.sigmoid(nd.dot(X, W_xf) + nd.dot(H, W_hf) + b_f)
        O = nd.sigmoid(nd.dot(X, W_xo) + nd.dot(H, W_ho) + b_o)
        C_tilda = nd.tanh(nd.dot(X, W_xc) + nd.dot(H, W_hc) + b_c)
        C = F * C + I * C_tilda
        H = O * C.tanh()
        Y = nd.dot(H, W_ho) + b_h
        outputs.append(Y)
    return outputs, (H, C)
```

The ABN ASIA logo, featuring a stylized blue feather or arrow pointing to the right, with the text 'ABN ASIA.ORG' below it.

Moving to Deep Learning

After mastering the basics, I needed to learn how to actually implement these concepts. I found several excellent resources that combined practical knowledge with theoretical understanding, including a particularly challenging but invaluable deep learning book.

Essential resources:

- Framework documentation
- Technical research papers
- Implementation guides
- Community tutorials

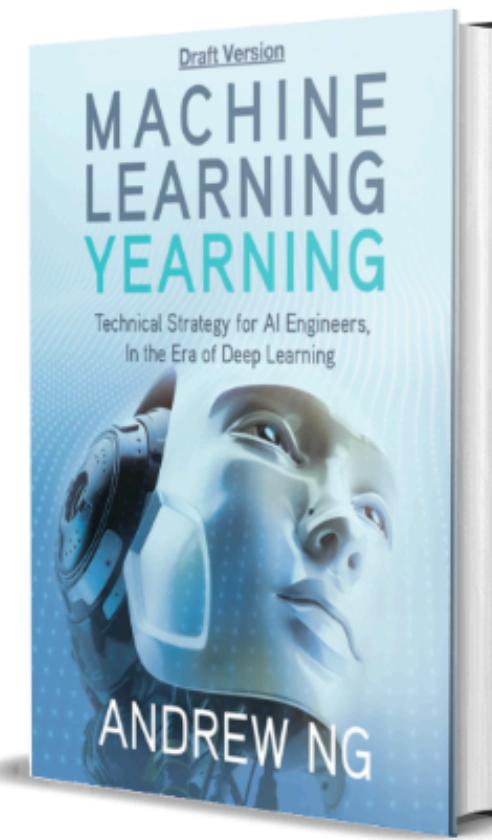
Resources

FREE EBOOK

Get The Machine Learning Yearning Book By Andrew NG

An introductory book about developing ML algorithms

[DOWNLOAD NOW](#)



Deep Learning

An MIT Press book

Ian Goodfellow and Yoshua Bengio and Aaron Courville

[Exercises](#) [Lectures](#) [External Links](#)

The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free.

The deep learning textbook can now be ordered on [Amazon](#).

For up to date announcements, join our [mailing list](#).

Citing the book

To cite this book, please use this bibtex entry:

```
@book{Goodfellow-et-al-2016,
  title={Deep Learning},
  author={Ian Goodfellow and Yoshua Bengio and Aaron Courville},
  publisher={MIT Press},
  note={\url{http://www.deeplearningbook.org}},
  year={2016}
}
```

To write your own document using our LaTeX style, math notation, or to copy our notation page, download our [template](#) files.



Hands-on Learning Method

My most effective learning technique was surprisingly old-school - manually retyping code from various GitHub repositories. While it might seem inefficient, this method forced me to understand every line of code and its purpose in the bigger picture.

Learning technique:

- Manual code rewriting
- Deep code analysis
- Pattern recognition
- Implementation practice

Resources

README MIT license

Zero to Mastery Deep Learning with TensorFlow

All of the course materials for the [Zero to Mastery Deep Learning with TensorFlow course](#).

This course will teach you the foundations of deep learning and how to build and train neural networks for various problem types with TensorFlow/Keras.

Important links

- Watch the [first 14-hours of the course on YouTube](#) (notebooks 00, 01, 02)
- Read the [beautiful online book version of the course](#)
- [Sign up](#) to the full course on the Zero to Mastery Academy (videos for notebooks 03-10)
- Got questions about the course? Check out the [livestream Q&A for the course launch](#)
- Get a quick overview of TensorFlow with the [TensorFlow Cheatsheet](#)

deep-learning-keras-tf-tutorial Public

Watch 19 ▾

master 1 Branch Tags Go to file t Add file ▾ Code ▾

dhavalsays	Update handling_imbalanced_data_exercise.md	0981fe5 · 3 years ago	31 Commits
10_gpu_benchmarking	tf mirrored stratergy	3 years ago	
11_chrun_prediction	ENH: moved all the active tutorial content from py/deeple...	3 years ago	
12_precision_recall	ENH: moved all the active tutorial content from py/deeple...	3 years ago	
13_dropout_layer	ENH: moved all the active tutorial content from py/deeple...	3 years ago	
14_imbalanced	Update handling_imbalanced_data_exercise.md	3 years ago	
16_cnn_cifar10_small_image_classification	cifar small image tutorial update	3 years ago	
17_data_augmentation	ENH: moved all the active tutorial content from py/deeple...	3 years ago	



Real-World Applications

The real learning began with actual problem-solving. I started with classic machine learning challenges and gradually moved to more complex problems. This hands-on experience was crucial for understanding how ML works in practice.

Project progression:

- Basic classification problems
- Advanced regression tasks
- Time series predictions
- Market analysis systems

Resources

KAGGLE · GETTING STARTED PREDICTION COMPETITION · ONGOING

Titanic - Machine Learning from Disaster

Start here! Predict survival on the Titanic and get familiar with ML basics

Overview Data Code Models Discussion Leaderboard Rules

HOME CREDIT GROUP · FEATURED PREDICTION COMPETITION · 6 YEARS AGO

Home Credit Default Risk

Can you predict how capable each applicant is of repaying a loan?

Overview Data Code Models Discussion Leaderboard Rules

KAGGLE · GETTING STARTED PREDICTION COMPETITION · ONGOING

House Prices - Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting

Overview Data Code Models Discussion Leaderboard Rules



Deployment Expertise

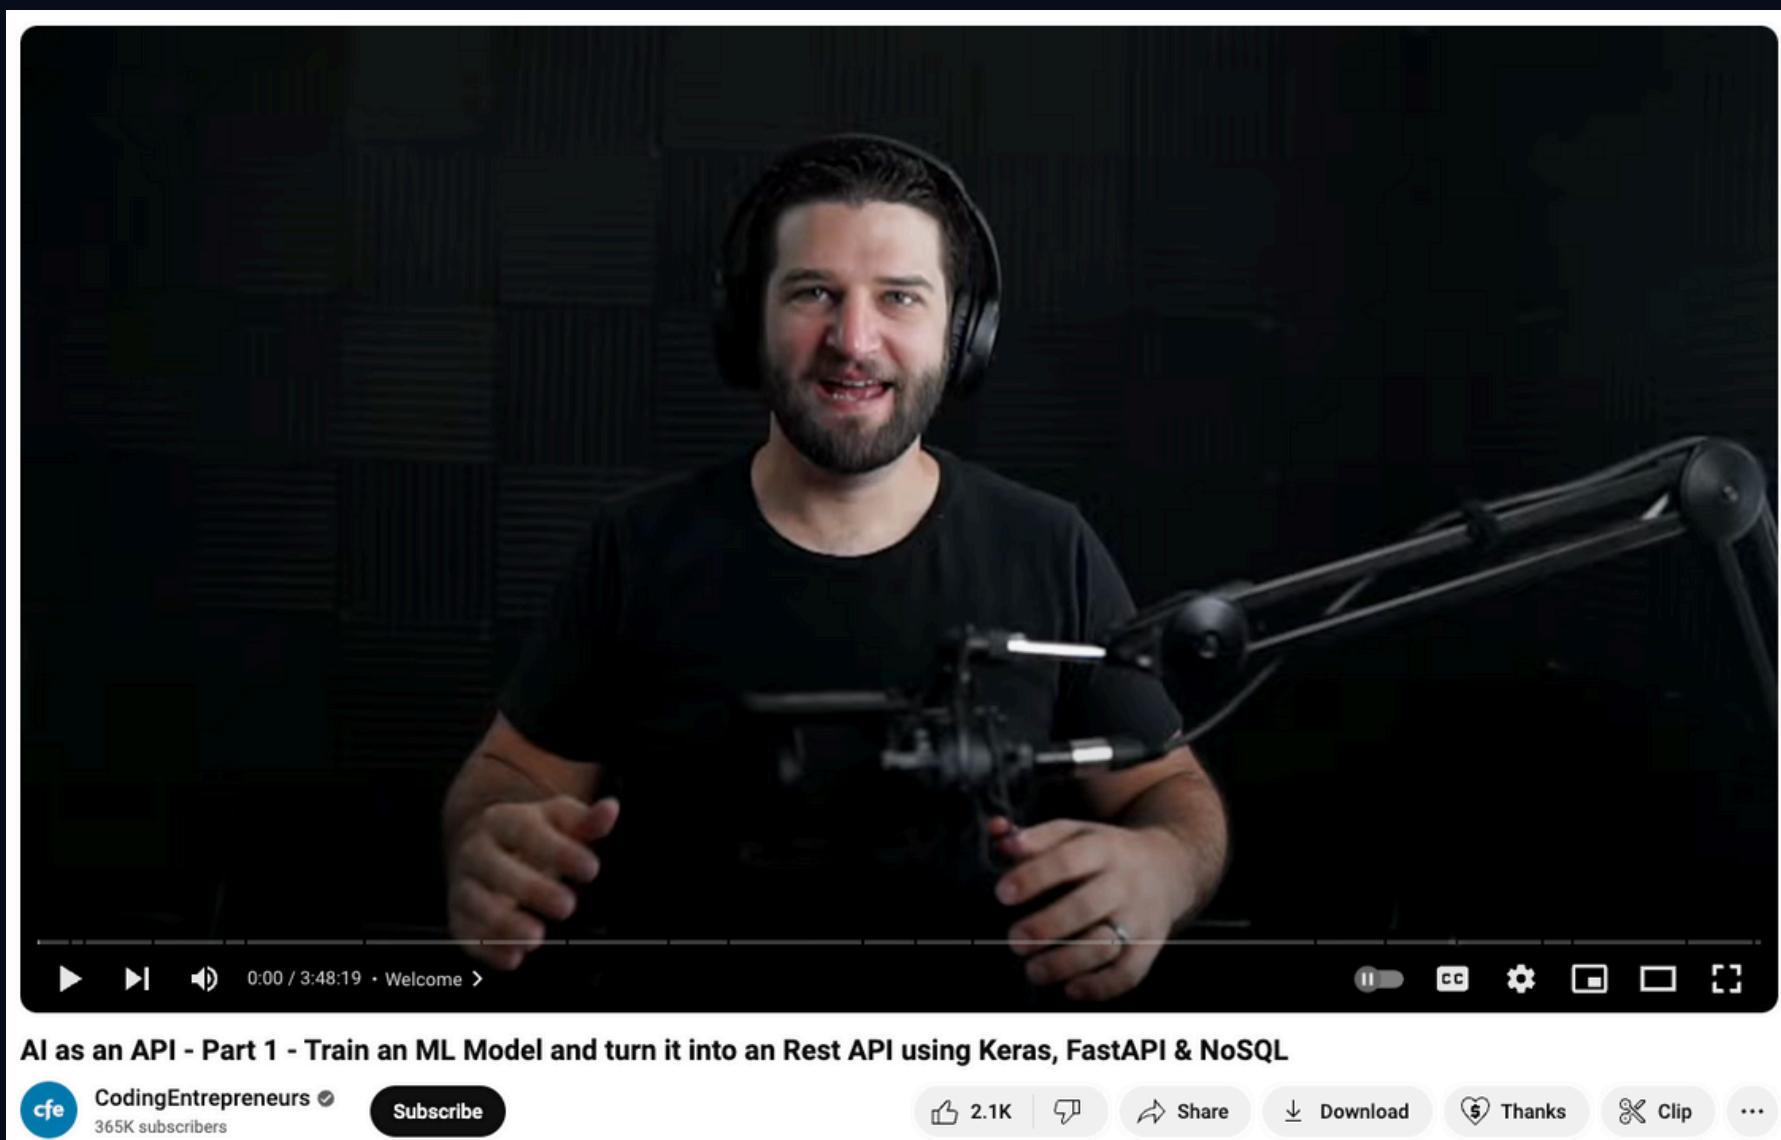
Before specializing further, I needed to learn how to deploy ML models effectively. I discovered that a simple framework and the right cloud provider can make deployment straightforward and cost-effective, whether you're working with basic models or complex GPT systems.

Deployment essentials:

- Fast API implementation
- Cloud service utilization
- Endpoint creation
- Cost optimization



Resources



A step-by-step approach to building a fast API for deep learning classification projects

Abdulwajeed Ganiyu · Follow
8 min read · Jul 4, 2023

43 2

```
31     self.file = None
32     self.fingerprints = set()
33     self.logdupes = True
34     self.debug = debug
35     self.logger = logging.getLogger(__name__)
36     if path:
37         self.file = open(os.path.join(path, 'fingerprint.log'), 'a')
38         self.file.seek(0)
39         self.fingerprints.update(self.file.read().split(os.linesep))
40
41     @classmethod
42     def from_settings(cls, settings):
43         debug = settings.getbool('DEBUG', False)
44         return cls(job_dir(settings), debug)
45
46     def request_seen(self, request):
47         fp = self.request_fingerprint(request)
48         if fp in self.fingerprints:
49             return True
50         self.fingerprints.add(fp)
51         if self.file:
52             self.file.write(fp + os.linesep)
53
54     def request_fingerprint(self, request):
```

Specializing in NLP (2020)

My career took a significant turn when I specialized in Natural Language Processing. Starting with the Transformer architecture and BERT models, I gradually moved to more advanced applications. This foundation is crucial as 99% of today's LLM models build on these concepts.

NLP progression:

- Transformer architecture
- Classification tasks
- GPT model adaptation
- Fine-tuning techniques



Resources

Attention Is All You Need

Welcome to the 😊 Course!



This course will teach you about natural language processing (NLP) using libraries from the [Hugging Face](#) ecosystem — 😊 [Transformers](#), 😊 [Datasets](#), 😊 [Tokenizers](#), and 😊 [Accelerate](#) — as well as the [Hugging Face Hub](#). It's completely free and without ads.

Illustrated Guide to Transformers- Step by Step Explanation



Michael Phi · Follow

Published in Towards Data Science · 12 min read · May 1, 2020

3.1K

21



...

Illustrated Guide to Transformers

Step by Step

I am fine <end>

Transformers

HI,
how

Transformers



ABN ASIA.ORG

Where I Am Today

After six years in ML, I've developed numerous projects, including an AI well-being support system and advanced stock market analysis tools. My doctoral dissertation combines ML, stock market prediction, and reinforcement learning, representing the culmination of this journey.

Current projects:

- AI well-being mobile app
- Stock market ML systems
- Doctoral research
- Open-source contributions



Show Case

sintelly

Home How it Works Blog Tutorials About Us

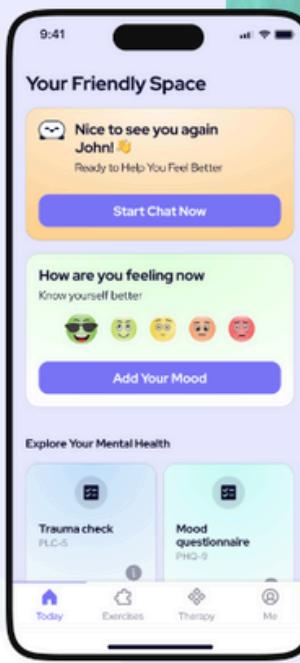
Your Trusted Ally for Mental Health Support

Experience Sintelly: Your Journey to Enhanced Mental Health Begins Now



9.69K reviews ★★★★★

4.5



PRIMO INVESTING

Home How it Works Features About FAQs

Unlock Free AI Insights for Better Investing

Get weekly, data-driven insights from our advanced Agent AI system.

Join the waitlist:

Join 128 investors already on the waitlist!

Key benefits:

- ✓ AI-driven stock analysis and strategies
- ✓ Market insights powered by GPT and RL
- ✓ Time-saving comprehensive market research
- ✓ Free weekly newsletter with expert AI analysis

